

Practical Work: Process Affinity

Master CSMI
Compilation & Performance
Bérenger Bramas

October 14, 2020

1 Summary

In this practical work, you will analyze how the pinning of the processes works and can be beneficial, but also what aligned memory allocation means.

2 Ressources

- `taskset` command man page: <https://linux.die.net/man/1/taskset>
- `sched_getcpu` man page: https://linux.die.net/man/3/sched_getcpu
- `sched_getaffinity` man page: https://linux.die.net/man/2/sched_getaffinity
- `aligned_alloc` man page: https://en.cppreference.com/w/c/memory/aligned_alloc

3 Practical work organization (always the same)

In the practical work, you will obtain the code from my repository and push it to your repository. Therefore, you will have to clone one branch per session and push it to your own repository. It is mandatory that you **commit and push** frequently (after each question and at the end of the session, at least) such that I can easily look at what you have coded at the end of the session, how you did progress (and potentially compare it with the latest version you will have).

It is required that you filled the `report.md` file to let me know what you did.

In the rest of the document, we consider you have a repository named `cnp-tp-2020` on `git.unistra.fr` that is private but that I can access in read.

3.1 Get the practical work

Consider you are in your project directory do the following:

```
# Clone my repo
git clone https://git.unistra.fr/bbramas/csmi-tp-2020.git --branch=TP2 csmi-tp2
# If you use SSH, use:
# git clone git@git.unistra.fr:bbramas/csmi-tp-2020.git --branch=TP2 csmi-tp2
# Go in the newly created directory
cd csmi-tp2
```

3.2 Add your repository as remote

You will push on your own repository:

```
# Rename my remote
git remote rename origin old-origin
# Add your own remote
git remote add origin https://git.unistra.fr/[YOU LOGIN HERE]/cnp-tp-2020.git
# If you use an SSH key:
```

```
# git remote add origin git@git.unistra.fr:[YOU LOGIN HERE]/cnp-tp-2020.git
# Push the current branch and active the tracking
git push -u origin TP2
```

3.3 During the session and while you work on the project

After each question or important modification push the current changes:

```
# No matter where you are in the project directory
git commit -a -m "I did something"
git push
```

3.4 When you are done

You have fully finished your work (at most D+14 H-2):

```
git commit -a -m "I did something"
git push
```

3.5 Important!

Do not forget that you have to fill the report.md file and commit it. Remember to commit regularly to keep track of your work and let me see a history of it if I need it. Do not share any code with someone else, as I am here to answer all questions and support all of you. Remember that you have questions to answer in the moodle before the end of the session and that you must push your branch at the end of the session too. Additional credits can be obtained if you make some modifications after the session to improve your solution. Changes can be made until two weeks after the session minus two hours, ie you must push before the beginning of the $n+2$ practical work. **Do not remove code from the test functions as I use them to evaluate your code.** Therefore, if you need you can add extra functions for your own testing/debugging. If some of them are showing interesting things about your code, simply leave a comment in the code and the report.md.

3.6 Compilation

To compile, we use CMake:

```
cd TP2
mkdir build
cd build
cmake ..
make # Will make all
make something # Will build only something
VERBOSE=1 make # Will show the commands used to compile (including the flags)
```

4 Reminder

4.1 Processor affinity

As we have seen in class, a CPU can have multiple cores, where each core can execute a different process. The OS is in charge of distributing the processes on the cores and decides based on criteria (priority, etc.) and heuristics which processes should be executed. Being preempted - remove from a core or move to another core - have significant execution penalty in terms of performance, because it means that the context of the process has to be stored and restored and that the data that were moved close to the core (in the different levels of the cache) might need to be reloaded from higher-levels. Therefore, when you execute simulations and try to get the best performance, or if you need to do an accurate execution time measure, the process must be pinned (bind) to a core. To pin a process, one has to inform the Linux

scheduler about it, and there are several ways to do so. The first approach is to inform the scheduler before an execution starts, it has the advantage of being transparent for the target application (this can be achieved with *taskset*). The second method is to use the programming interface to the Linux scheduler to express the fact that a given process should run on some specific core(s). The good news is that this can be done directly from the process itself (see *sched_getaffinity*).

4.2 Aligned memory

Memory address p is say to be X -aligned if $p \bmod X$ is 0. In current systems only power of two alignment are meaningful, and p is aligned to 2^Y if the first non zero bit in p is at position Y . For instance, 1101.0101 is aligned to 1, but 1010.1000 is aligned to 8.

By default the classical allocation functions do not perform any specific alignment, and thus the addresses they return should be considered as aligned to 1 even if they might be aligned to a higher degree. One of the reasons is that the regular allocators focus more and filling the blank (having efficient memory filling pattern) rather than returning aligned blocks. Having aligned memory matters when we deal with performance because the caches work with lines of 64 Bytes (that must be 64 aligned). If a value at position $@v$ is used, the cache line that is 64 aligned and which includes v will be moved $((@v/64) * 64)$. Moreover, if the compiler knows that a block is aligned, it can use specific instructions to load them into register, as we will see when we will vectorize a code.

There are different possibilities to have aligned allocations. The first one is to use existing functions, such as the C11 *aligned_alloc*. The second is to create your own allocator on top of *malloc/new*.

5 Pinning

5.1 Print the available CPU core

The file *print_av_cores.cpp* contains the code to know the pinning of the current process, and print the list on the standard output. Using the *taskset* command try to do the following:

- pin the process on a single core
- pin the process on two cores

You should obtain output like:

```
$ ./print_av_cores
Available cores =
0 1 2 3 4 5 6 7

$ taskset [something] ./print_av_cores
Available cores =
0

$ taskset [something] ./print_av_cores
Available cores =
0 1
```

5.2 Develop a kernel to print the used cores

Update the *print_moves.cpp* file to print everytime the process is moved on a different core. To ensure it does, run several instance of the application on multiple terminals. To know on which core is the current process, you can use the function *sched_getcpu*.

Example of desired output:

```
$ ./print_moves
>> Move to core 3
>> Core count : [0] 1    [3] 1
>> Move to core 7
>> Core count : [0] 1    [3] 1    [7] 1
>> Move to core 3
```

```

>> Core count : [0] 1    [3] 2    [7] 1
>> Move to core 7
>> Core count : [0] 1    [3] 2    [7] 2
>> Move to core 3
>> Core count : [0] 1    [3] 3    [7] 2
>> Move to core 7
>> Core count : [0] 1    [3] 3    [7] 3
>> Move to core 3
>> Core count : [0] 1    [3] 4    [7] 3
>> Move to core 7
>> Core count : [0] 1    [3] 4    [7] 4

```

5.3 Measure the performance cost of changing cores

Update the file `dot_with_move.cpp` such that in the second scope the process is moved on a different core after each call to `dot`. Notice that a list of available cores is obtained at the beginning of the main.

Example of desired output:

```

$ ./dot_with_move
Check dot
TestSize = 500000
>> Without move : 1.22568
>> With move : 4.4557

```

6 Memory alignment

6.1 Memory alignment evaluation

Fill the function `alignementOfPtr` in the first part of the code in file `aligned_malloc.cpp` such that the main will print the alignment of several imaginary addresses but also the addresses returned by C11 aligned malloc function. You can consider only alignment power of 2.

```

Test with hard numbers
Address 1
>> Alignement 1
Address 2
>> Alignement 2
Address 4
>> Alignement 4
Address 8
>> Alignement 8
Address 6
>> Alignement 2
Address 7
>> Alignement 1
Test with C11
alignment = 1
Address 0x5585a4a86280
>> Alignement 128
...
alignment = 2
Address 0x5585a4a86280
>> Alignement 128
...
alignment = 4
Address 0x5585a4a86280
>> Alignement 128
...
alignment = 8

```

```

Address 0x5585a4a86280
>> Alignement 128
...
alignment = 16
Address 0x5585a4a86280
>> Alignement 128
...

```

6.2 Create an allocator for aligned memory

Create your own function that will allocate aligned memory. It is advised to do it in two steps. In the first one, you can simply allocate more memory than asked, and return a pointer that is correctly aligned (but that will not be compatible with the regular free leading to memory leaks). In a second step, you will store the address of the real block just before the aligned pointer and use it to free as usual.

```

// First step
function aligned-malloc(alignment, size)
    size_with_extra_space = extra_need_space(alignment, size)
    unaligned_ptr = malloc(size)
    ptr = get_first_aligned_address_in_block(alignment, unaligned_ptr)
    return ptr
end

// Second step
function aligned-malloc(alignment, size)
    size_with_extra_space = extra_need_space(alignment, size)
    unaligned_ptr = malloc(size)
    ptr = get_first_aligned_address_in_block(alignment, unaligned_ptr)
    *(ptr-size_of_pointer) = unaligned_ptr
    return ptr
end

function aligned-free(ptr)
    unaligned_ptr = *(ptr-size_of_pointer)
    free(unaligned_ptr)
end

```